

Annotation Guidelines for:

*Beyond Good Intentions:
When Does the Framing of Multilingual and Low-Resource NLP Research
Become a Caricature?*
(EMNLP2026)

If you use these guidelines, please cite:

*Beyond Good Intentions: When Does the Framing of Multilingual and
Low-Resource NLP Research Become a Caricature?*

Ousidhoum et al. [2026]

We examine the motivations, background assumptions, epistemic elements, and research framing (i.e., narratives) found in NLP papers that focus on multilingual resources and models, particularly those targeting underserved communities.

We adapt the two-step annotation scheme defined by Schlichtkrull et al. [2023] and Chamoun et al. [2025] to our goals: (1) paragraph-level annotation and (2) discourse-level narrative annotation. The paragraph-level annotation identifies passages related to the papers’ goals and methodology. From these excerpts, we determine the means (i.e., applications and models) and ends of the papers. We then extract the implied narratives and assess whether these narratives realistically (or plausibly) achieve the intended aims. Finally, we identify several forms of epistemic harm that may result from the ways in which resources are depicted (e.g., saviourism).

1 Paper Selection

We collect 100 research papers that focus on multilingual NLP, and NLP for low-resource languages in particular. We examine each paper’s introduction and abstract to extract quotes.

First, we annotated a small set of papers based on the criteria defined below. Then, we discussed our annotations and decided on the final set of aspects to label and their respective categories. Finally, we discussed all annotations, unified a set of slightly confusing labels that caused disagreements, and annotated the papers for a final round based on the following criteria.

1.1 Paper Metadata

The type of the paper survey, dataset, model description, task description, other (name).

The subfield/task translation, NER, instruction tuning, ...

The language(s) Number of languages and list if any.

2 Annotation Criteria

We analyse the introduction and abstract of each paper and define epistemic narratives at two levels: paragraph and discourse.

Accordingly, each paper is annotated in a two-step process: (1) extracting elements at the paragraph level, and (2) inferring the corresponding narratives at the discourse level.

2.1 Paragraph-Level Annotation

At the paragraph level, we extract epistemic elements related to the paper’s means, ends, and evidential support. Specifically, we annotate means (i.e., applications), ends (i.e., goals), and evidential support used to justify the claimed contribution. We also extract relevant quotes and identify the actors and subjects involved where these are relevant to the paper’s narrative framing. At the discourse level, we combine these elements into narratives—e.g., “*researchers* (actors) should use our *dataset* (means) to *train language models* (means) created by *volunteers* (data workers) who are speakers from region X (data subjects), in order to *serve speakers of the continent* (end).”

We work inductively, starting with a set of predefined categories based on an initial pilot. If, during the annotation process, we encounter a means, end, evidential-support type, or narrative that does not fit any of the predefined labels, we introduce a new label. We use a multi-label approach, as individual paragraphs may contain multiple references to categories.

Important Note In this annotation stage, we do not extract implied statements or infer relationships between elements. That is, we rely on extracted quotes and do not interpret them. Interpretation and inference are reserved for the discourse-level narrative annotation.

2.1.1 Quotes

We extract the relevant quote or paragraph itself. We focus on epistemic quotes, i.e., passages related to knowledge, justification, motivation, or the claimed contribution. We organise these quotes around questions relevant to the paper’s narrative framing, such as the *why* and *what* of the research. We initially extract

quotes from the introduction; however, if relevant information is missing there, we consult the abstract and extract it when available.

2.1.2 Evidential Support

We extract any evidence or justification used to support claims made in the paper’s framing. This may include:

- **Non-NLP scientific articles:** evidence drawn from research in other scientific domains;
- **Newspaper articles or polls:** media reports or collected public opinion;
- **Prior NLP work:** evidence drawn from previous NLP papers;
- **Anecdotes:** isolated examples, personal observations, or individual experiences used as justification;
- **Expressions of threat:** claims about potential harm or negative consequences that are not substantiated by evidence;
- **Opinions:** strong assertions or claims presented without citation or empirical support.

For example, ‘NLP has been shown to revive languages through building resources [CITE a *CL* paper]’ would be annotated as *Prior NLP work*, whereas a benchmark for [minority] language X does not exist because of ethnic tensions” would be annotated as an *Opinion* if no supporting evidence is provided. Similarly, “if we do not build an LLM for [minority] language X, it will disappear” would be annotated as an *Expression of threat*.

2.1.3 Means/Application

Means refer to how authors propose to use what is developed in the paper to accomplish a particular goal or end. We identify the following categories:

- **Improving LLM performance.** Developing resources, methods, or models intended to improve the performance of LLMs for particular languages or tasks. E.g., *Most LLMs perform poorly in multilingual contexts. Our collected instructions improve their performance for language X across tasks A, B, and C.*
- **NLP applications.** Developing or improving technologies for specific tasks, such as machine translation, named entity recognition, transcription, and other applications. E.g., *Translation to language X is poor. Our model focuses on translation for n languages.*
- **LLM safety and ethical considerations.** Developing or applying methods intended to make language models safer or more appropriate for particular linguistic communities. E.g., *We aim to make LLMs safer for speakers of language X.*

- **Improving LLM evaluation.** Developing methods, benchmarks, or metrics for evaluating and comparing the capabilities of LLMs, particularly multilingual models, and identifying their limitations.
- **Helping language learners.** Developing NLP technologies to support people learning one or more languages. E.g., *By building better technologies for language X, we help people who are trying to learn it.*
- **Other.** Applications that do not fit within the categories above, such as data collection or annotation.

2.1.4 Ends

Ends refer to the purpose or ultimate aim of the work—that is, what the authors claim to accomplish. We identify the following categories:

- **Serving communities.** Goals explicitly directed towards benefiting a specific group of people, for example by improving NLP technologies for a particular linguistic community.
- **Reducing inequality.** Broader goals of improving representation, coverage, or equality beyond a specific community. This includes generalising an observed gap from one language or language family to multiple groups that do not necessarily belong to the same community. In contrast to *serving communities*, this category does not explicitly focus on a specific group.
- **Advancing knowledge production.** Contributing new insights, resources, or methods that support scientific research (i.e., helping researchers).
- **Improving [LLM] evaluation.** Developing tools, benchmarks, or metrics that support the evaluation of NLP systems/LLMs.
- **Language revitalisation.** Supporting the preservation, maintenance, or revitalisation of a language through data collection and NLP technologies.
- **Other.** Goals that do not fit the categories above, including addressing environmental issues (e.g., reducing the carbon footprint of NLP systems through greater efficiency), building powerful models or achieving state-of-the-art performance, and government monitoring where the stated aim involves monitoring government interests or activities in a particular region.

2.2 Discourse-Level Annotation—Narratives

At the discourse level, we identify the epistemic narratives present in each paper. We define epistemic narratives as framings based on the relationships among the stated means, ends, and evidential support. We identify the following narrative categories:

- **Knowledge Production.** A narrative in which the primary justification is scientific contribution. The stated goal is to advance understanding or support future research without linking the work to broader societal outcomes.
- **Explicit Pathway to Impact.** A narrative in which both means and ends are clearly articulated. The paper describes a concrete technical contribution (e.g., an automated language-processing task) and links it to a defined outcome (e.g., serving a specific community). This category excludes cases where the sole stated end is knowledge production.
- **Vague Serving.** A narrative in which a technical contribution is presented as broadly serving a language community or population without a clearly articulated pathway to impact or supporting evidence.
- **Saviourism / Decolonisation.** A narrative in which a resource or model is presented as saving an endangered language, reviving a historical language, or contributing to decolonisation.
- **Tech Solutionism.** A narrative in which a technical resource or tool is presented as addressing inherently social, political, or structural problems (e.g., regional tensions) without supporting evidence or a demonstrated mechanism, potentially exaggerating the impact of the intervention.
- **Tokenism.** A narrative in which limited acts of inclusion (e.g., involving annotators from underrepresented groups) are presented as meaningfully addressing broader structural inequalities in ways that may simplify or exaggerate the social impact of the work. This narrative is more common in promotional materials.

References

- Eric Chamoun, Nedjma Ousidhoum, Michael Schlichtkrull, and Andreas Vlachos. Social good or scientific curiosity? uncovering the research framing behind nlp artefacts. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 25310–25346, 2025.
- Nedjma Ousidhoum, Noopur Zambare, and Mohamed Abdalla. Beyond good intentions: When does the framing of multilingual and low-resource nlp research become a caricature? In *Proceedings of EMNLP*, 2026.
- Michael Schlichtkrull, Nedjma Ousidhoum, and Andreas Vlachos. The intended uses of automated fact-checking artefacts: Why, how and who. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 8618–8642, 2023.